

## Load Libraries

```
In [58]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import Preprocessing
from sklearn.preprocessing import LabelEncoder, OneHotEncoder, StandardScaler
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import r2_score
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
%matplotlib inline
sns.set()
```

```
In [79]: df = pd.read_csv( 'car_price_prediction.csv')
df
```

Out[79]:

	ID	Price	Levy	Manufacturer	Model	Prod. year	Category	Leather interior	Fu ty
0	45654403	13328	1399	LEXUS	RX 450	2010	Jeep	Yes	Hybr
1	44731507	16621	1018	CHEVROLET	Equinox	2011	Jeep	No	Petr
2	45774419	8467	-	HONDA	FIT	2006	Hatchback	No	Petr
3	45769185	3607	862	FORD	Escape	2011	Jeep	Yes	Hybr
4	45809263	11726	446	HONDA	FIT	2014	Hatchback	Yes	Petr
...	...	...	...	...	...	...	...	...	...
19232	45798355	8467	-	MERCEDES-BENZ	CLK 200	1999	Coupe	Yes	CN
19233	45778856	15681	831	HYUNDAI	Sonata	2011	Sedan	Yes	Petr
19234	45804997	26108	836	HYUNDAI	Tucson	2010	Jeep	Yes	Dies
19235	45793526	5331	1288	CHEVROLET	Captiva	2007	Jeep	Yes	Dies
19236	45813273	470	753	HYUNDAI	Sonata	2012	Sedan	Yes	Hybr

19237 rows × 18 columns



In [80]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 19237 entries, 0 to 19236
Data columns (total 18 columns):
#   Column                Non-Null Count  Dtype
---  -
0   ID                    19237 non-null  int64
1   Price                 19237 non-null  int64
2   Levy                  19237 non-null  object
3   Manufacturer          19237 non-null  object
4   Model                 19237 non-null  object
5   Prod. year            19237 non-null  int64
6   Category              19237 non-null  object
7   Leather interior      19237 non-null  object
8   Fuel type             19237 non-null  object
9   Engine volume         19237 non-null  object
10  Mileage               19237 non-null  object
11  Cylinders              19237 non-null  float64
12  Gear box type         19237 non-null  object
13  Drive wheels          19237 non-null  object
14  Doors                 19237 non-null  object
15  Wheel                 19237 non-null  object
16  Color                 19237 non-null  object
17  Airbags               19237 non-null  int64
dtypes: float64(1), int64(4), object(13)
memory usage: 2.6+ MB
```

## Clean Dataset

```
In [81]: '''
Preprocessing:
Drop Duplicates - Replacing categorical values - Clean Outliers - Add a new Feat
'''
df = Preprocessing.preprocessing_pipeline(df)
```

```
Preprocessing started...
Initial shape: (19237, 18)
After dropping duplicates: (18924, 18)
Replacing categorical values...
After cleaning outliers: (16035, 18)
Feature engineering...
Dropping columns...
Final shape: (16035, 16)
```

```
In [82]: # Handling the Object columns and Transform to numerical value
df['Levy'] = df['Levy'].replace('-', '0')
df['Levy'] = pd.to_numeric(df['Levy'])

df['Engine volume'] = df['Engine volume'].replace('Turbo', '')
df['Engine volume'] = pd.to_numeric(df['Engine volume'])

df['Mileage km'] = df['Mileage'].replace('km', '',)
df['Mileage km'] = pd.to_numeric(df['Mileage km'])

df.drop(columns='Mileage', inplace=True)
```

```
In [ ]:
```

# EDA:

In [6]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Index: 16035 entries, 0 to 18921
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Price                 16035 non-null  int64
1   Levy                 16035 non-null  int64
2   Manufacturer         16035 non-null  object
3   Model                16035 non-null  object
4   Category             16035 non-null  object
5   Leather interior     16035 non-null  object
6   Fuel type            16035 non-null  object
7   Engine volume        16035 non-null  float64
8   Cylinders            16035 non-null  float64
9   Gear box type        16035 non-null  object
10  Drive wheels         16035 non-null  object
11  Wheel                16035 non-null  object
12  Color                16035 non-null  object
13  Airbags              16035 non-null  int64
14  Age                  16035 non-null  int64
15  Mileage km           16035 non-null  int64
dtypes: float64(2), int64(5), object(9)
memory usage: 2.1+ MB
```

In [7]: `df.isna().sum()`

```
Out[7]: Price                0
Levy                      0
Manufacturer              0
Model                    0
Category                 0
Leather interior         0
Fuel type                0
Engine volume            0
Cylinders                 0
Gear box type            0
Drive wheels             0
Wheel                    0
Color                    0
Airbags                  0
Age                      0
Mileage km               0
dtype: int64
```

In [8]: `df.describe().round(2).T`

Out[8]:

	count	mean	std	min	25%	50%	75%	max
<b>Price</b>	16035.0	14294.82	11287.26	1.0	5217.0	12544.0	20385.0	47120.0
<b>Levy</b>	16035.0	575.38	457.74	0.0	0.0	640.0	862.0	2209.0
<b>Engine volume</b>	16035.0	2.13	0.60	0.8	1.6	2.0	2.5	3.5
<b>Cylinders</b>	16035.0	4.37	0.88	1.0	4.0	4.0	4.0	16.0
<b>Airbags</b>	16035.0	6.57	4.25	0.0	4.0	6.0	12.0	16.0
<b>Age</b>	16035.0	14.22	5.52	5.0	11.0	13.0	16.0	86.0
<b>Mileage km</b>	16035.0	130490.58	80334.73	0.0	70994.0	124892.0	180000.0	363661.0

In [9]: `df.select_dtypes(include='object').describe()`

Out[9]:

	Manufacturer	Model	Category	Leather interior	Fuel type	Gear box type	Drive wheels	Wheel	C
<b>count</b>	16035	16035	16035	16035	16035	16035	16035	16035	16035
<b>unique</b>	59	1318	11	2	7	4	3	2	2
<b>top</b>	TOYOTA	Prius	Sedan	Yes	Petrol	Automatic	Front	Left wheel	B
<b>freq</b>	3232	989	7514	11174	8162	11208	11525	14668	3

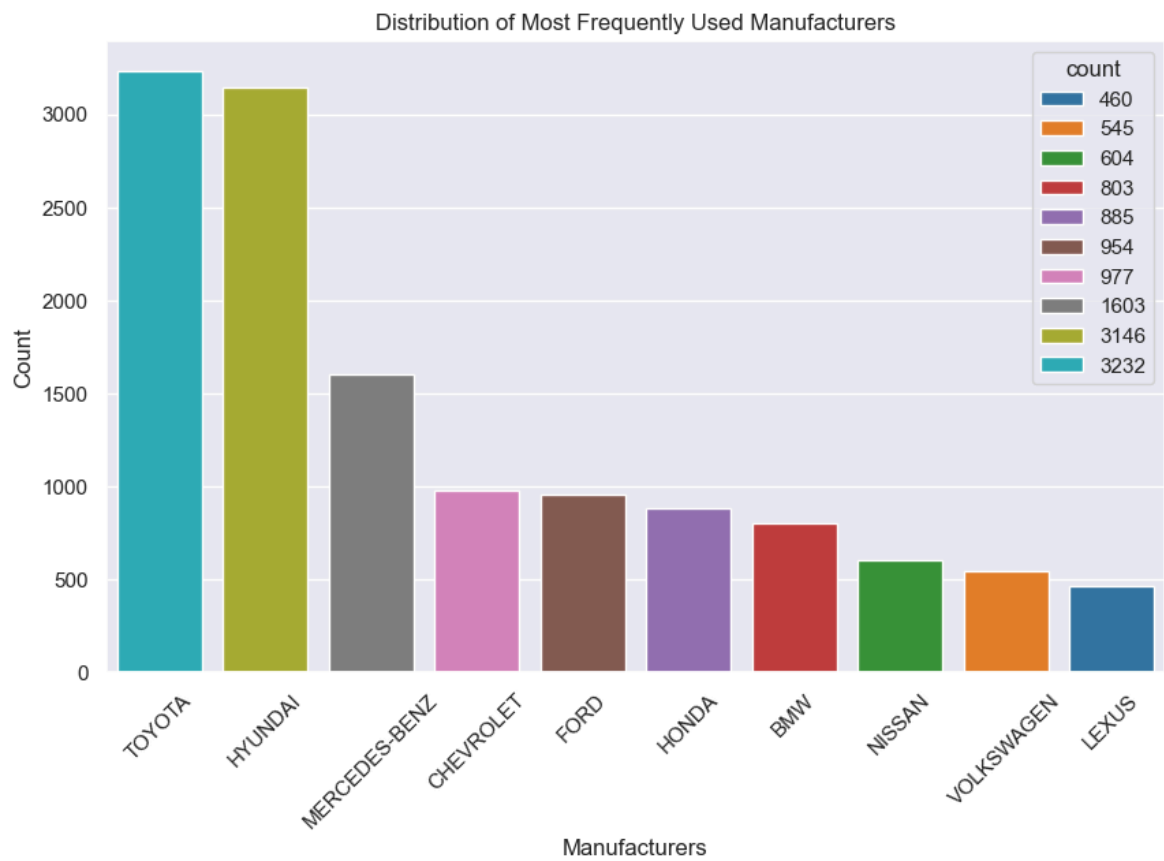
In [10]: `# Most frequently used manufacturer  
df['Manufacturer'].value_counts().head(10)`

Out[10]: Manufacturer

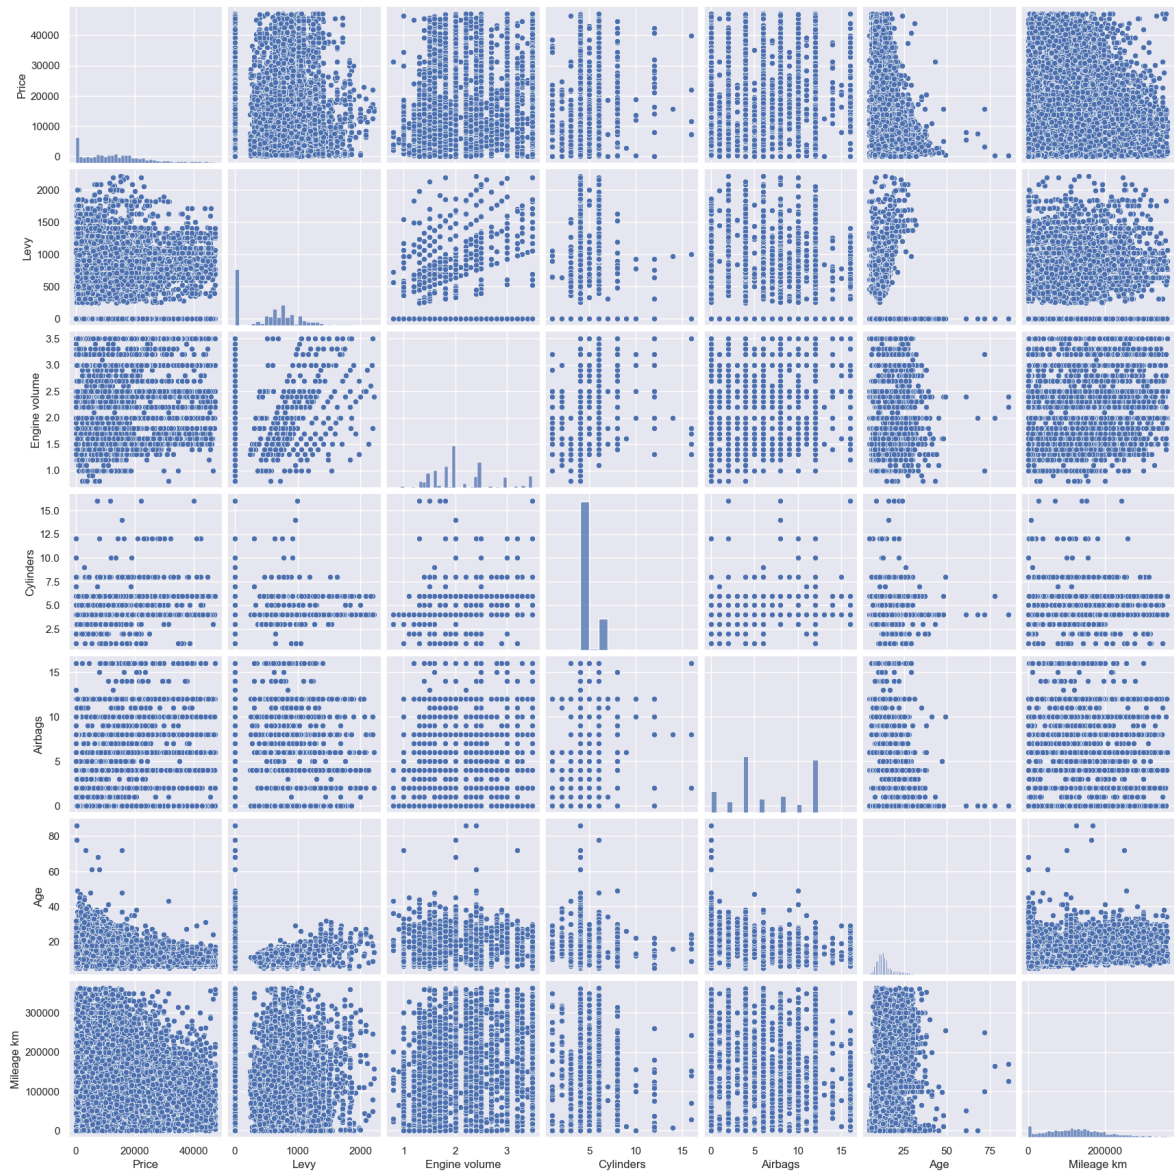
TOYOTA	3232
HYUNDAI	3146
MERCEDES-BENZ	1603
CHEVROLET	977
FORD	954
HONDA	885
BMW	803
NISSAN	604
VOLKSWAGEN	545
LEXUS	460

Name: count, dtype: int64

In [11]: `# Top 10 most frequently used manufacturers  
top_manufacturers = df['Manufacturer'].value_counts().head(10)  
  
plt.figure(figsize=(10,6))  
sns.barplot(x= top_manufacturers.index, y=top_manufacturers.values,hue=top_manuf  
plt.title('Distribution of Most Frequently Used Manufacturers')  
plt.xticks(rotation=45)  
plt.xlabel('Manufacturers')  
plt.ylabel('Count')  
plt.show()`



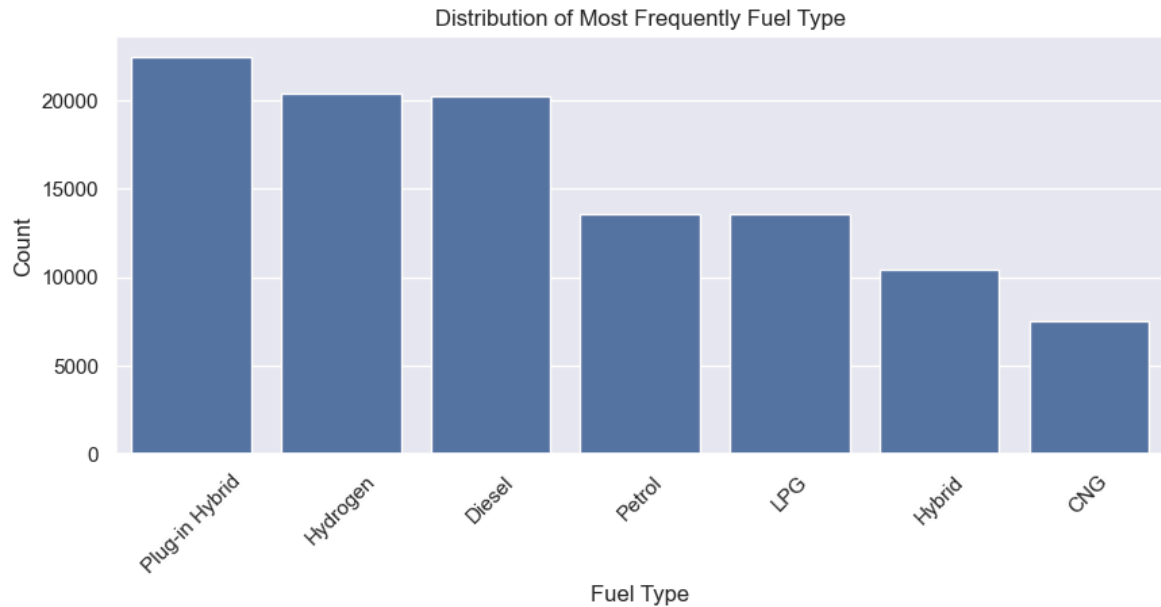
```
In [12]: #Relationship between the columns  
sns.pairplot(df)  
plt.show()
```



```
In [13]: # most frequently fuel type
fuel_type = df.groupby('Fuel type')['Price'].mean().round(1).sort_values(ascending=True)
fuel_type
```

```
Out[13]: Fuel type
Plug-in Hybrid    22445.0
Hydrogen          20385.0
Diesel            20245.0
Petrol            13580.9
LPG               13547.9
Hybrid            10463.2
CNG               7539.1
Name: Price, dtype: float64
```

```
In [14]: plt.figure(figsize=(10,4))
sns.barplot(x= fuel_type.index, y=fuel_type.values)
plt.title('Distribution of Most Frequently Fuel Type')
plt.xticks(rotation=45)
plt.xlabel('Fuel Type')
plt.ylabel('Count')
plt.show()
```



```
In [15]: #Relationship between numerical columns
numeric_column = df.select_dtypes('number')
numeric_column.corr().round(2)
```

```
Out[15]:
```


	Price	Levy	Engine volume	Cylinders	Airbags	Age	Mileage km
Price	1.00	0.02	0.00	-0.05	-0.07	-0.30	-0.17
Levy	0.02	1.00	0.31	0.11	0.21	-0.46	-0.04
Engine volume	0.00	0.31	1.00	0.65	0.31	0.02	0.19
Cylinders	-0.05	0.11	0.65	1.00	0.22	0.12	0.15
Airbags	-0.07	0.21	0.31	0.22	1.00	-0.27	-0.01
Age	-0.30	-0.46	0.02	0.12	-0.27	1.00	0.23
Mileage km	-0.17	-0.04	0.19	0.15	-0.01	0.23	1.00

```
In [16]: #Columns that affect Price
price_corr = numeric_column.corr()['Price'].sort_values(ascending=False)
price_corr
```

```
Out[16]: Price      1.000000
Levy        0.019057
Engine volume 0.003895
Cylinders   -0.052854
Airbags     -0.068340
Mileage km  -0.171228
Age         -0.296334
Name: Price, dtype: float64
```

```
In [17]: plt.figure(figsize=(10,6))
sns.heatmap(numeric_column.corr().round(2), annot=True, cmap= 'rainbow')
plt.title('The Relationship Between The Numerical Columns')
plt.show()
```






```
In [18]: #Count of cars that have genuine Leather interior and that have artificial leather interior
leather_interior = df.groupby('Leather interior')['Price'].count()
leather_interior
```

```
Out[18]: Leather interior
No      4861
Yes     11174
Name: Price, dtype: int64
```


```
In [19]: plt.figure(figsize=(5,3))
sns.countplot(df, x='Leather interior', hue='Leather interior', palette='coolwarp')
plt.title('Count of Cars that have Leather interior or No')
plt.show()
```



```
In [20]: #Most used gearbox
gear_box = df.groupby('Gear box type')['Price'].count()
gear_box
```

```
Out[20]: Gear box type
Automatic    11208
Manual       1668
Tiptronic    2451
Variator      708
Name: Price, dtype: int64
```

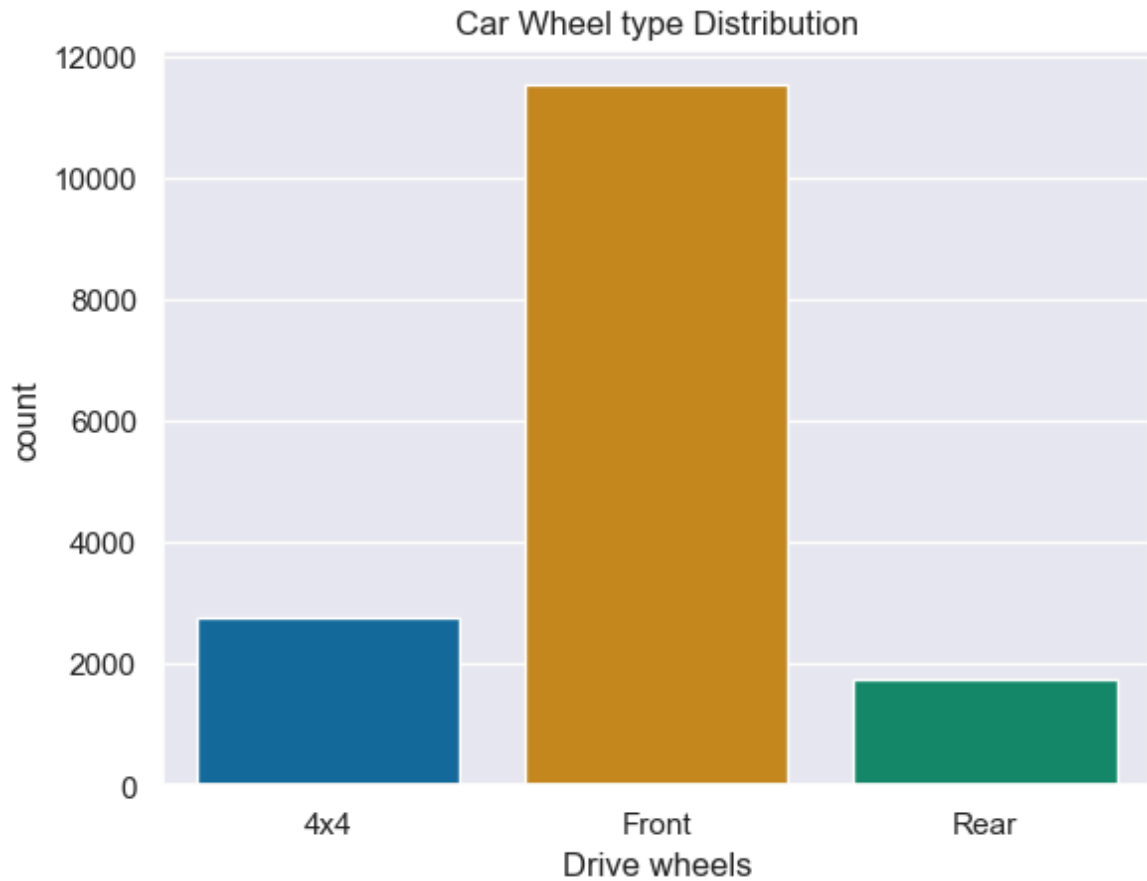
```
In [21]: sns.countplot(df,x='Gear box type',hue='Gear box type', palette= 'rainbow')
plt.title('The Most Commonly Used Type of Gearbox')
plt.show()
```



```
In [22]: #Most used Drive Wheels type
drive_wheels= df.groupby('Drive wheels')['Price'].count()
drive_wheels
```

```
Out[22]: Drive wheels
4x4       2753
Front     11525
Rear       1757
Name: Price, dtype: int64
```

```
In [23]: sns.countplot(df,x= 'Drive wheels',hue='Drive wheels', palette= 'colorblind')
plt.title('Car Wheel type Distribution')
plt.show()
```



## Machine Learning:

In [30]: `df.head()`

Out[30]:

	Price	Levy	Manufacturer	Model	Category	Fuel type	Engine volume	Cylinders	Color	Airb
0	13328	1399	28	1035	4	2	3.5	6.0	12	
1	16621	1018	6	563	4	5	3.0	6.0	1	
2	8467	0	18	585	3	5	1.3	4.0	1	
3	3607	862	13	565	4	2	2.5	4.0	14	
4	11726	446	18	585	3	5	1.3	4.0	12	

5 rows × 23 columns



```
In [83]: # Define the columns to be one-hot encoded
one_hot_columns = ['Leather interior', 'Gear box type', 'Drive wheels', 'Wheel']

# Create a OneHotEncoder object with the desired settings
oh_encoder = OneHotEncoder(sparse_output=False, handle_unknown='ignore')

# Fit the encoder and transform the training data
oh_encoder_train = oh_encoder.fit_transform(df[one_hot_columns])

# Get the new column names after one-hot encoding
```

```

oh_encoded_columns = oh_encoder.get_feature_names_out(one_hot_columns)

# Create a new DataFrame for the one-hot encoded columns
oh_encoded_train_df = pd.DataFrame(oh_encoder_train, columns=oh_encoded_columns,

# Concatenate the original DataFrame with the one-hot encoded DataFrame
df = pd.concat([df, oh_encoded_train_df], axis=1)

# Drop the original columns that were one-hot encoded
df.drop(columns=one_hot_columns, inplace=True)

```

```

In [84]: # Define the columns to be Label encoded
label_encode_columns = ['Manufacturer', 'Model', 'Category', 'Fuel type', 'Color

# Create a dictionary to store the Label encoders for each column
label_encoders = {}

# Iterate over each column to apply Label encoding
for column in label_encode_columns:
    # Create a LabelEncoder object
    label_encoder = LabelEncoder()

    # Fit the encoder and transform the column in the DataFrame
    df[column] = label_encoder.fit_transform(df[column])

    # Store the fitted label encoder in the dictionary
    label_encoders[column] = label_encoder

```

```
In [ ]:
```

## 1- Linear Regression

```

In [85]: lr = LinearRegression()
scores_linear_regression = cross_val_score(lr, X, y, cv=5, scoring='r2')
print(f'R2 Score with Linear Regression = {scores_linear_regression.mean().round

```

R2 Score with Linear Regression = 0.24

## 2- Random Forest Regressor

```

In [86]: rf = RandomForestRegressor()
scores_random_forest = cross_val_score(rf, X, y, cv=5, scoring='r2')
print(f'R2 Score with Random Forest Regressor = {scores_random_forest.mean().rou

```

```

C:\Users\student7\anaconda3\Lib\site-packages\sklearn\base.py:1473: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
    return fit_method(estimator, *args, **kwargs)
C:\Users\student7\anaconda3\Lib\site-packages\sklearn\base.py:1473: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
    return fit_method(estimator, *args, **kwargs)
C:\Users\student7\anaconda3\Lib\site-packages\sklearn\base.py:1473: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
    return fit_method(estimator, *args, **kwargs)
C:\Users\student7\anaconda3\Lib\site-packages\sklearn\base.py:1473: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
    return fit_method(estimator, *args, **kwargs)
C:\Users\student7\anaconda3\Lib\site-packages\sklearn\base.py:1473: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
    return fit_method(estimator, *args, **kwargs)
R2 Score with Random Forest Regressor = 0.77

```

In [ ]:

## Deep Learning:

In [87]: `X = StandardScaler().fit_transform(X)`

```

In [88]: ann = Sequential()
ann.add(Dense(25, input_dim= 22, activation='relu')) #input

ann.add(Dense(625, activation='relu')) #hidden

ann.add(Dense(1,activation='linear')) #output

```

```

C:\Users\student7\anaconda3\Lib\site-packages\keras\src\layers\core\dense.py:87:
UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.
    super().__init__(activity_regularizer=activity_regularizer, **kwargs)

```

In [89]: `ann.summary()`

Model: "sequential\_3"

Layer (type)	Output Shape
dense_9 (Dense)	(None, 25)
dense_10 (Dense)	(None, 625)
dense_11 (Dense)	(None, 1)



Total params: 17,451 (68.17 KB)

Trainable params: 17,451 (68.17 KB)

**Non-trainable params: 0 (0.00 B)**

```
In [90]: ann.compile(optimizer='adam', loss='mean_squared_error')
```

```
In [91]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_
```

```
In [92]: hist= ann.fit(X_train,y_train, epochs=100, batch_size=25, verbose=1)
```

```
Epoch 1/100
514/514 ————— 1s 2ms/step - loss: 304765504.0000
Epoch 2/100
514/514 ————— 1s 2ms/step - loss: 99093088.0000
Epoch 3/100
514/514 ————— 1s 2ms/step - loss: 86885568.0000
Epoch 4/100
514/514 ————— 1s 1ms/step - loss: 84314024.0000
Epoch 5/100
514/514 ————— 1s 1ms/step - loss: 84023512.0000
Epoch 6/100
514/514 ————— 1s 2ms/step - loss: 80020856.0000
Epoch 7/100
514/514 ————— 1s 2ms/step - loss: 78018256.0000
Epoch 8/100
514/514 ————— 1s 2ms/step - loss: 79333424.0000
Epoch 9/100
514/514 ————— 1s 2ms/step - loss: 78017736.0000
Epoch 10/100
514/514 ————— 1s 2ms/step - loss: 75039536.0000
Epoch 11/100
514/514 ————— 1s 2ms/step - loss: 75010448.0000
Epoch 12/100
514/514 ————— 1s 2ms/step - loss: 75903280.0000
Epoch 13/100
514/514 ————— 1s 2ms/step - loss: 72387528.0000
Epoch 14/100
514/514 ————— 1s 2ms/step - loss: 72255848.0000
Epoch 15/100
514/514 ————— 1s 2ms/step - loss: 73274344.0000
Epoch 16/100
514/514 ————— 1s 2ms/step - loss: 70961760.0000
Epoch 17/100
514/514 ————— 1s 2ms/step - loss: 69953408.0000
Epoch 18/100
514/514 ————— 1s 2ms/step - loss: 67956784.0000
Epoch 19/100
514/514 ————— 1s 2ms/step - loss: 68356096.0000
Epoch 20/100
514/514 ————— 1s 2ms/step - loss: 66821184.0000
Epoch 21/100
514/514 ————— 1s 2ms/step - loss: 65494604.0000
Epoch 22/100
514/514 ————— 1s 2ms/step - loss: 63735976.0000
Epoch 23/100
514/514 ————— 1s 2ms/step - loss: 62734204.0000
Epoch 24/100
514/514 ————— 1s 2ms/step - loss: 61502452.0000
Epoch 25/100
514/514 ————— 1s 2ms/step - loss: 61591064.0000
Epoch 26/100
514/514 ————— 1s 2ms/step - loss: 59354372.0000
Epoch 27/100
514/514 ————— 1s 2ms/step - loss: 57166568.0000
Epoch 28/100
514/514 ————— 1s 2ms/step - loss: 54785160.0000
Epoch 29/100
514/514 ————— 1s 2ms/step - loss: 54924224.0000
Epoch 30/100
514/514 ————— 1s 2ms/step - loss: 54371556.0000
```

```
Epoch 31/100
514/514 ————— 1s 2ms/step - loss: 50525308.0000
Epoch 32/100
514/514 ————— 1s 2ms/step - loss: 50854356.0000
Epoch 33/100
514/514 ————— 1s 2ms/step - loss: 50095600.0000
Epoch 34/100
514/514 ————— 1s 2ms/step - loss: 49564228.0000
Epoch 35/100
514/514 ————— 1s 2ms/step - loss: 47634500.0000
Epoch 36/100
514/514 ————— 1s 2ms/step - loss: 48587860.0000
Epoch 37/100
514/514 ————— 1s 2ms/step - loss: 46469396.0000
Epoch 38/100
514/514 ————— 1s 2ms/step - loss: 48332208.0000
Epoch 39/100
514/514 ————— 1s 2ms/step - loss: 48324068.0000
Epoch 40/100
514/514 ————— 1s 2ms/step - loss: 47665772.0000
Epoch 41/100
514/514 ————— 1s 2ms/step - loss: 46887376.0000
Epoch 42/100
514/514 ————— 1s 2ms/step - loss: 46719536.0000
Epoch 43/100
514/514 ————— 1s 2ms/step - loss: 46522112.0000
Epoch 44/100
514/514 ————— 1s 2ms/step - loss: 44215292.0000
Epoch 45/100
514/514 ————— 1s 2ms/step - loss: 45187036.0000
Epoch 46/100
514/514 ————— 1s 2ms/step - loss: 45744004.0000
Epoch 47/100
514/514 ————— 1s 2ms/step - loss: 43647804.0000
Epoch 48/100
514/514 ————— 1s 2ms/step - loss: 45254292.0000
Epoch 49/100
514/514 ————— 1s 1ms/step - loss: 46476420.0000
Epoch 50/100
514/514 ————— 1s 2ms/step - loss: 45242312.0000
Epoch 51/100
514/514 ————— 1s 2ms/step - loss: 45000052.0000
Epoch 52/100
514/514 ————— 1s 2ms/step - loss: 43651512.0000
Epoch 53/100
514/514 ————— 1s 2ms/step - loss: 45041604.0000
Epoch 54/100
514/514 ————— 1s 2ms/step - loss: 45271284.0000
Epoch 55/100
514/514 ————— 1s 2ms/step - loss: 44024032.0000
Epoch 56/100
514/514 ————— 1s 2ms/step - loss: 45567520.0000
Epoch 57/100
514/514 ————— 1s 2ms/step - loss: 45127076.0000
Epoch 58/100
514/514 ————— 1s 2ms/step - loss: 44853144.0000
Epoch 59/100
514/514 ————— 1s 2ms/step - loss: 44972316.0000
Epoch 60/100
514/514 ————— 1s 2ms/step - loss: 44729316.0000
```



```
Epoch 61/100
514/514 ————— 1s 2ms/step - loss: 43066952.0000
Epoch 62/100
514/514 ————— 1s 2ms/step - loss: 43586984.0000
Epoch 63/100
514/514 ————— 1s 2ms/step - loss: 45392864.0000
Epoch 64/100
514/514 ————— 1s 2ms/step - loss: 45365012.0000
Epoch 65/100
514/514 ————— 1s 2ms/step - loss: 43374176.0000
Epoch 66/100
514/514 ————— 1s 2ms/step - loss: 43206288.0000
Epoch 67/100
514/514 ————— 1s 2ms/step - loss: 45235908.0000
Epoch 68/100
514/514 ————— 1s 2ms/step - loss: 43240308.0000
Epoch 69/100
514/514 ————— 1s 2ms/step - loss: 43073660.0000
Epoch 70/100
514/514 ————— 1s 2ms/step - loss: 43547672.0000
Epoch 71/100
514/514 ————— 1s 2ms/step - loss: 43944820.0000
Epoch 72/100
514/514 ————— 1s 2ms/step - loss: 42863504.0000
Epoch 73/100
514/514 ————— 1s 2ms/step - loss: 43050120.0000
Epoch 74/100
514/514 ————— 1s 2ms/step - loss: 44079348.0000
Epoch 75/100
514/514 ————— 1s 2ms/step - loss: 41865744.0000
Epoch 76/100
514/514 ————— 1s 2ms/step - loss: 44269896.0000
Epoch 77/100
514/514 ————— 1s 2ms/step - loss: 42601840.0000
Epoch 78/100
514/514 ————— 1s 2ms/step - loss: 42896972.0000
Epoch 79/100
514/514 ————— 1s 2ms/step - loss: 42076788.0000
Epoch 80/100
514/514 ————— 1s 2ms/step - loss: 42522552.0000
Epoch 81/100
514/514 ————— 1s 2ms/step - loss: 43827352.0000
Epoch 82/100
514/514 ————— 1s 2ms/step - loss: 42613732.0000
Epoch 83/100
514/514 ————— 1s 2ms/step - loss: 44798056.0000
Epoch 84/100
514/514 ————— 1s 2ms/step - loss: 42482340.0000
Epoch 85/100
514/514 ————— 1s 2ms/step - loss: 44200848.0000
Epoch 86/100
514/514 ————— 1s 2ms/step - loss: 42278164.0000
Epoch 87/100
514/514 ————— 1s 2ms/step - loss: 42647160.0000
Epoch 88/100
514/514 ————— 1s 2ms/step - loss: 41852220.0000
Epoch 89/100
514/514 ————— 1s 2ms/step - loss: 43746044.0000
Epoch 90/100
514/514 ————— 1s 2ms/step - loss: 42720916.0000
```

```

Epoch 91/100
514/514 ————— 1s 2ms/step - loss: 44095084.0000
Epoch 92/100
514/514 ————— 1s 2ms/step - loss: 42318800.0000
Epoch 93/100
514/514 ————— 1s 2ms/step - loss: 41620384.0000
Epoch 94/100
514/514 ————— 1s 2ms/step - loss: 44715224.0000
Epoch 95/100
514/514 ————— 1s 2ms/step - loss: 42595912.0000
Epoch 96/100
514/514 ————— 1s 2ms/step - loss: 43198044.0000
Epoch 97/100
514/514 ————— 1s 2ms/step - loss: 41303592.0000
Epoch 98/100
514/514 ————— 1s 2ms/step - loss: 43372816.0000
Epoch 99/100
514/514 ————— 1s 2ms/step - loss: 43033144.0000
Epoch 100/100
514/514 ————— 1s 2ms/step - loss: 43761436.0000

```

```
In [93]: hist.history.keys()
```

```
Out[93]: dict_keys(['loss'])
```

```
In [94]: y_pred= ann.predict(X_test)
score_ann = r2_score(y_test,y_pred)
print(f'R2 Score with ANN = {score_ann:.2f}')
```


```

101/101 ————— 0s 1ms/step
R2 Score with ANN = 0.64

```

## R2 score Comparison:

```
In [101... plt.figure(figsize=(8,4))
models = ['Linear Regression', 'Random Forest', 'ANN']
scores = [0.24, 0.77, 0.66]
plt.bar(models,scores, color='skyblue')
plt.title('R2 Score Comparison')
plt.ylabel('R2 Score')
plt.show()
```



## Summary

- Random Forest performed best with  $R^2 = 0.77$
- The performance of the ANN was okay, and could be improved
- Linear Regression underperformed, likely due to data complexity
- Futuer improvement: hyperparameter tuning, deep model enhancements

## Thank You